SafeMAP: Analysis-Guided Planning and Per-Function Validation for Safe C-to-Rust Translation

Jacob Sebastian Cyril[1*], Benita Jose Chalissery[1] and NewBegin M[1]

[1*]Department of Computer Science, CHRIST (Deemed to be) University, Bangalore, India.

*Corresponding author(s). E-mail(s): jacob.sebastian@mca.christuniversity.in;
Contributing authors: benita.jaison@christuniversity.in; new.begin@christuniversity.in;

Abstract

Automatic C-to-Rust translators either keep C's pointer conventions in unsafe Rust or rely on a large language model (LLM) whose output is hard to trust. SafeMAP translates a C function only when static analysis finds a safe, testable Rust interface for it, and refuses the rest with a recorded reason. Planned functions are synthesized by exact rules or a plan-guided LLM request with compiler-feedback repair. A translation is accepted only if it contains no `unsafe`, exposes no raw pointers, and matches the original C function on up to 1,000 generated inputs, with the C side run under AddressSanitizer and UndefinedBehaviorSanitizer. On ten held-out C libraries (356 translation units) selected under a sealed protocol, SafeMAP accepted 12–16 units across three runs. The same LLM accepted 9–15 when given SafeMAP's planned interfaces and 1–2 when given only the C sources; raw C2Rust accepted none. Most units were refused before translation because they need manual refactoring, such as removing global state (159), or use C types with no safe interface (116). The acceptance check rejected 91% of non-equivalent compiling mutants of accepted functions. Ablations attribute the accepted translations to interface planning rather than to prompt guidance or repair. A post-hoc analysis shows that per-unit lint checking would raise the held-out result to 39 units. Per-function, oracle-checked acceptance is thus feasible on production C but narrow, and the recorded refusals show what currently blocks safe automatic migration.

Keywords: C-to-Rust migration, safe Rust, static analysis, large language models, differential testing, memory safety

1 Introduction

C remains common in systems software. Its memory model leaves programmers responsible for preventing buffer overflows, use-after-free errors, null dereferences, and data races. Safe Rust checks ownership and borrowing at compile time while retaining low-level control, so migrating C code to safe Rust is attractive. Translation is difficult because many C idioms have no direct safe-Rust equivalent: raw pointers, output parameters, manual allocation, nullable values, and integer error codes all require changes in representation and interface.

C2Rust preserves C behavior by generating Rust that uses raw pointers and unsafe operations (Immunant n.d.). Later work reduces this unsafety by static analysis, ownership inference, refactoring, and LLM rewriting (Emre et al. 2021; Hong 2023; Zhang et al. 2023; Ling et al. 2022; Nitin et al. 2026; Mohapatra et al. 2026). A second line asks LLMs to translate directly to safe Rust and checks the result with tests, fuzzing, or verification (Hong and Ryu 2025b; Zhou et al. 2025a; Eniser et al. 2025; Shetty et al. 2024; Zhou et al. 2025b; Yang et al. 2025; Sim et al. 2026; Shiraishi et al. 2026; Farrukh et al. 2026; Li et al. 2025). These systems generally aim to translate everything and then measure how much of the result is safe or correct.

SafeMAP takes the opposite position. It translates a function only when static analysis finds a safe Rust interface for it, and it accepts a translation only when that function, on its own, passes a behavioral check against the original C. Everything else is refused with a recorded reason. The question is therefore not how much unsafe code remains, but which translated functions can be trusted: how many functions of real C code can be migrated to safe Rust with per-function evidence of equivalence, and what prevents the rest. Our results locate SafeMAP's contribution in the planning step: an LLM told the planned safe interface produces checkable translations, whereas the same LLM given only the C code almost never does.

An earlier, frozen evaluation of SafeMAP on three C libraries accepted no units: defects in interface planning, rule matching, and harness construction each forced the outcome to zero (Section 4.1). This paper reports the corrected system, evaluated under a new sealed protocol on ten libraries that were not used during development. We ask:

- **RQ1:** How many translation units of production C libraries does SafeMAP accept as safe, behaviorally validated Rust?
- **RQ2:** Why are the remaining units not accepted?
- **RQ3:** How does SafeMAP compare with a one-shot LLM, with and without SafeMAP's planned interfaces, with its own rules alone, and with raw C2Rust, under one acceptance policy?
- **RQ4:** Which planning components contribute to acceptance?

- **RQ5:** How reliable is the acceptance decision, in terms of the harness's ability to reject wrong code and the variation across LLM runs?

This paper contributes:

- a translation pipeline that decides, per function and before any translation, whether a safe and testable Rust interface exists, and records a reason when it does not;
- a per-function acceptance procedure that compares each Rust function with the original C function under sanitizers, excluding inputs on which the C code has undefined behavior, and whose strength we measure by mutation analysis;
- a sealed, frozen evaluation on ten held-out libraries showing that planned interfaces raise a one-shot LLM from 1–2 to 9–15 accepted units, with repeated LLM runs, mutation analysis of the acceptance check, and component ablations; and
- a breakdown of 356 held-out units by the reason they could not be migrated, together with a public artifact for every decision.

2 Background and related work

2.1 C2Rust and safer-Rust refactoring

C2Rust translates C through a mostly syntactic process and is often used as a baseline or a migration starting point (Immunant n.d.). It preserves many C constructs with raw pointers and unsafe operations, so much of its output falls outside Rust's safe subset. Emre et al. found that raw pointers cause much of this unsafety (Emre et al. 2021) and later showed that Rust's borrow checker rejects many pointer uses that are in fact safe, which limits how far ownership inference can go (Emre et al. 2023). Studies of hand-written Rust find that even experienced developers use `unsafe` widely, mostly for interoperation and aliasing data structures (Astrauskas et al. 2020). Hong used static analysis to replace selected unsafe constructs in C2Rust output, including lock APIs and output parameters (Hong 2023). C2SaferRust combines C2Rust, static analysis, and LLM rewriting to reduce unsafe code in translated real-world programs (Nitin et al. 2026). Crown infers ownership models for C pointers and converts many of them to safe Rust types (Zhang et al. 2023). CRustS and Cp2SRust use different source-to-source transformations and transpiler designs (Ling et al. 2022; Mohapatra et al. 2026). Industry studies and surveys report that automatic translation still needs substantial manual work and testing (Hong and Ryu 2025a; Viirola et al. 2025).

2.2 LLM-based C-to-Rust translation

LLMs can produce idiomatic Rust, although correctness and type migration remain open problems. Hong and Ryu studied an LLM translator supplied with candidate signatures, caller and callee context, and compiler feedback (Hong and Ryu 2025b). C2RustTV checks LLM translations with generated tests and repair loops (Zhou et al. 2025a). Eniser et al. evaluate LLM translation through cross-language differential fuzzing and find lower success on more complex real-world samples (Eniser et al. 2025). RustFlow studies prompting and iterative repair (Zhang et al. 2026). Valenzuela et al. examine compiler errors in direct C-to-safe-Rust translation (Valenzuela et al. 2025).

RustMap combines program analysis and LLM support at project scale (Cai et al. 2025). Pointer knowledge graphs annotate code dependency graphs with ownership, mutability, and nullability facts to guide LLM translation of whole projects (Yuan et al. 2026), and His2Trans mines interfaces and idioms from earlier migrations of the same codebase (Wang et al. 2026c). Both supply the LLM with interface knowledge, which our results identify as the decisive input.

Several recent systems validate LLM translations against the C program. SmartC2Rust and SafeTrans repair translations iteratively from compiler errors and observed behavioral differences (Shiraishi et al. 2026; Farrukh et al. 2026). ACToR pairs a translator agent with an agent that builds differential fuzzers between the C and Rust binaries (Li et al. 2025), and TRAIL refines translation insights against executable counterexamples (Wang et al. 2026b). Syzygy translates units in dependency order and tests each one against its C counterpart with LLM-generated inputs (Shetty et al. 2024). SACTOR first produces interface-preserving Rust that may use `unsafe`, then refines it to idiomatic safe Rust, and checks both stages with end-to-end tests linked to the C code through the foreign-function interface (Zhou et al. 2025b). VERT checks a readable LLM translation against a WebAssembly-derived oracle with bounded model checking or property-based testing (Yang et al. 2025). ENCRUST wraps LLM translations in ABI-preserving shims and removes the remaining unsafe code with an agent (Sim et al. 2026), and Reboot translates whole interpreters to safe Rust through feature milestones with occasional user intervention (Wang et al. 2026a). CRUST-Bench supplies 100 C repositories with hand-written safe-Rust interfaces and tests (Khatry et al. 2025); those interfaces play the role that SafeMAP's planned signatures play here.

Table 1 places SafeMAP among these systems. SafeMAP differs in three ways. It decides per function, before translation, whether a safe interface exists, and records a reason when it does not instead of emitting `unsafe` code. Its oracle is the original C function itself, run under ASan and UBSan, and inputs on which C has undefined behavior are excluded rather than used as expected outputs. And acceptance is per function, so one failing unit does not hide the others. The cost of this conservatism is coverage, which Section 5 reports.

2.3 Undefined behavior and the C oracle

Differential testing compares two implementations on the same inputs and treats any disagreement as a defect (McKeeman 1998). It needs a trustworthy oracle. A C execution with undefined behavior does not define an expected output. ASan and UBSan detect many such executions (Serebryany et al. 2012; LLVM Project n.d.); SafeMAP runs every C oracle execution under both and discards inputs on which either reports an error. Miri can detect undefined behavior in Rust (Jung et al. 2026), but generated code that contains no `unsafe` cannot exhibit it, so SafeMAP does not run Miri.

Table 1 Validation and scope of recent C-to-Rust systems, as described by their authors.

System	Output	Behavioral oracle	Code it cannot translate safely
C2Rust (Immunant n.d.)	unsafe Rust	none	translated unsafely
Syzygy (Shetty et al. 2024)	safe Rust	C, LLM-generated inputs	repaired by LLM
SACTOR (Zhou et al. 2025b)	safe Rust	end-to-end tests via FFI	unsafe first stage
VERT (Yang et al. 2025)	safe Rust	Wasm-derived Rust oracle	regenerated
ENCRUST (Sim et al. 2026)	safe Rust	project test vectors	shims, then agent
SafeMAP	safe Rust	C under ASan/UBSan, per function	refused with reason

3 SafeMAP design

3.1 Acceptance policy

A translation unit is a C function, or a group of mutually recursive functions. SafeMAP accepts a unit if and only if all of the following hold:

1. it was admitted to the crate: it has the planned signature, contains no `unsafe`, raw pointers, `extern` blocks, or placeholders, and compiles together with the units already admitted;
2. the final crate compiles under `#![forbid(unsafe_code)]`, and its tests and Clippy pass;
3. its public interface exposes no raw pointer types; and
4. its per-unit differential test passes (Section 3.4).

A skipped or inapplicable check never counts as a pass. A unit whose callee was not admitted is recorded as `dependency_missing`.

3.2 Pipeline

Figure 1 shows the pipeline. For each C project, SafeMAP:

1. recovers compile information and analyzes every function with libclang: signatures, pointer roles (read, write, index, length pairing, nullability), allocation, calls, and unsupported constructs;
2. groups mutually dependent functions into units and orders them by dependency;
3. screens each unit (`candidate_safe`, `manual_refactor_required`, `unsafe_required`, `unknown`); mutable global state makes a unit `manual_refactor_required`;
4. plans a safe Rust signature for each candidate. The planner reports every C type it cannot map to a safe, testable Rust type: opaque `void *` or `FILE *`, pointer-to-pointer, returned aliasing pointers, and structs whose pointer fields have unresolved

ownership. A unit with any such issue is refused at planning; no signature is fabricated. A struct pointer may be planned as a Rust struct of the struct's scalar fields when every use of the pointer in the unit is a scalar-field access, a null test, or an argument to a function that satisfies the same condition;

5. synthesizes each planned unit in dependency order, by an exact deterministic rule when one applies and otherwise by a plan-guided LLM request followed by up to three compiler-feedback repairs; and
6. validates the crate and each unit against the acceptance policy.

3.3 Deterministic rules and plan-guided synthesis

Deterministic rules cover scalar expressions, boolean and bit-reversal idioms, pointer-length and fixed-size arrays, scalar structs and field access, dependency-ordered internal calls, output parameters, return-code plus output-parameter to `Result`, nullable pointers, reductions, simple C string lengths, and simple allocation-to-`Vec` and `Box` patterns. A rule that renders a single expression applies only when the whole function body is that statement and the expression references only parameters and called functions.

When no rule applies, SafeMAP sends the LLM the C function, the macros, enums, and constant tables it references, the Rust definitions of the planned structs, the signatures of already-accepted callees, the analysis facts, and the migration plan, and asks for exactly the planned signature. The response is admitted only if it passes the policy check and compiles with the units already admitted; otherwise the compiler diagnostics are returned for up to three repairs.

Consider the LLVM development helper below. libclang finds two indexed pointers paired with `n` and no writes, so the plan maps each pointer-length pair to a shared slice, and the `slice_dot_product` rule applies:

```
double loop(float *x, float *y, long n) {
  double acc = 0.0;
  for (long i=0; i<n; ++i) acc += (double)x[i] * (double)y[i];
  return acc;
}

pub fn r#loop(x: &[f32], y: &[f32]) -> f64 {
  x.iter().zip(y.iter()).map(|(&a,&b)| (a as f64)*(b as f64))
   .fold(0.0_f64, |s,v| s + v)
}
```

The initial value `+0.0` is necessary: differential testing showed that `sum()` can disagree with C on the sign of zero.

3.4 Per-unit differential harness

For each admitted unit, SafeMAP generates inputs from a fixed seed, including integer literals mined from the C body and their neighbors, and embeds the same literal tables in a C harness and a Rust harness. Slice and array arguments are exact-size heap buffers, so an out-of-bounds read in C is caught by ASan rather than masked. The C harness includes the original translation unit and runs under ASan and UBSan with

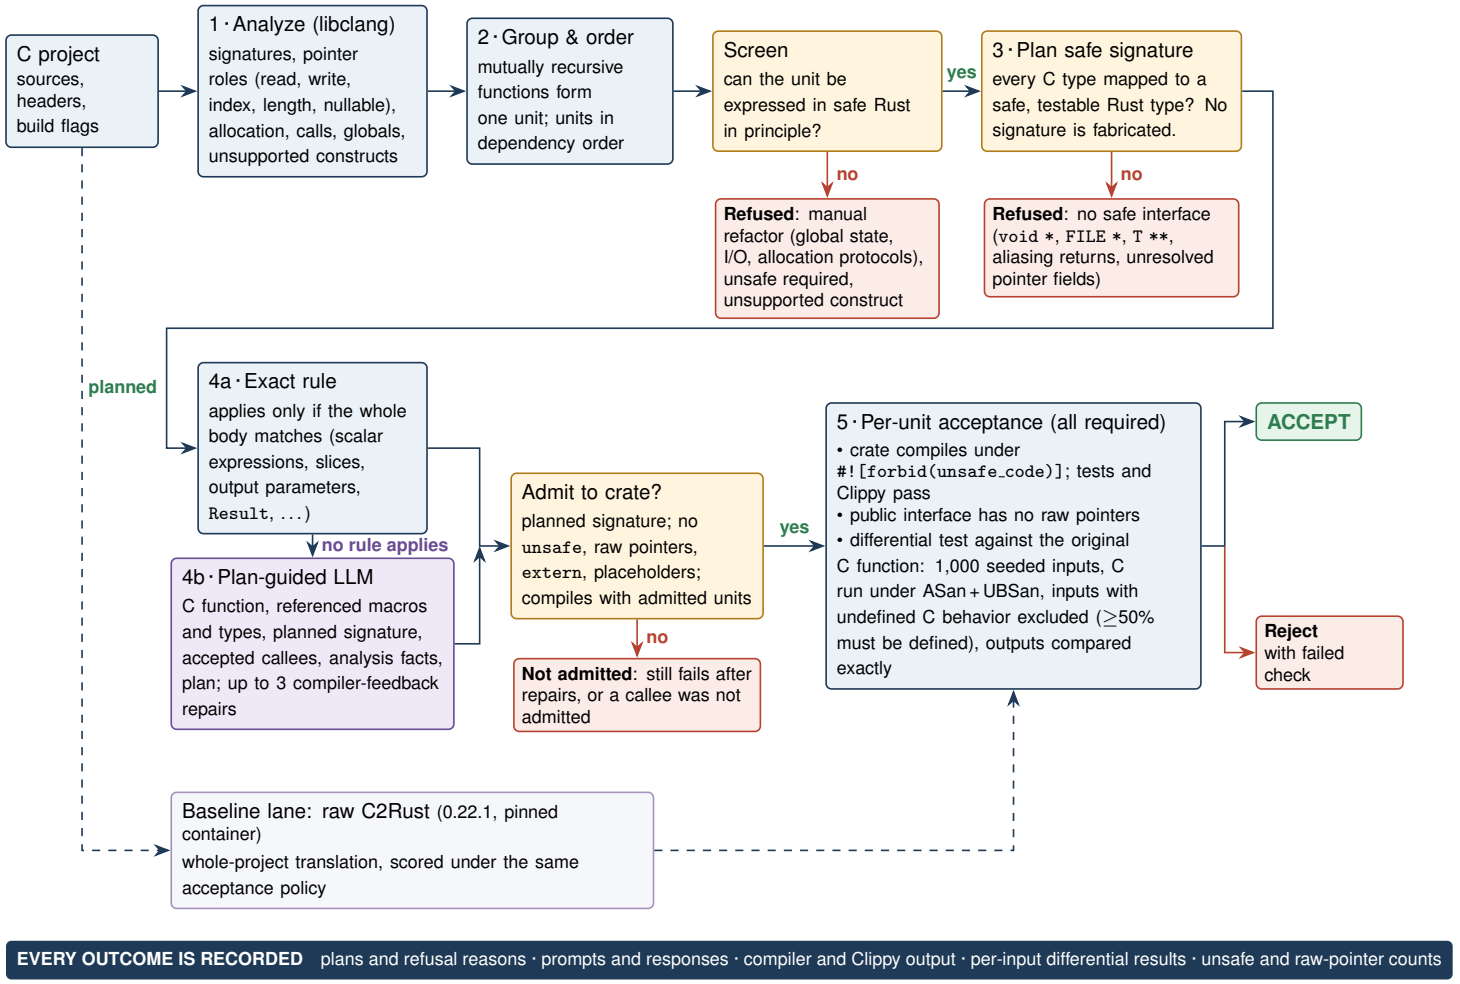


Fig. 1 SafeMAP pipeline. Units are refused at screening or at signature planning when no safe, testable interface exists. A planned unit is synthesized by an exact rule or, when none applies, by a plan-guided LLM request with compiler-feedback repair, and is accepted only if it passes every per-unit check, including a differential test against the original C function under sanitizers. Raw C2Rust is scored under the same policy. Every refusal and failed check is recorded with its reason

`-fno-sanitize-recover=all`; the Rust harness is built with overflow checks. Return values and every value written through a mutable parameter are compared exactly (floating point by bit pattern, strings byte by byte). An input on which the C oracle stops with a sanitizer report has no defined behavior and is excluded and counted. A defined C `abort` or `exit` must be matched by a Rust panic. A unit needs at least 50% defined inputs; otherwise it is recorded as `c_oracle_undefined_behavior`.

3.5 Baselines

Rules only runs the same pipeline without the LLM. *Raw C2Rust* translates each project in a pinned container (C2Rust 0.22.1, LLVM 14) and is measured under the same policy. The *one-shot LLM* baseline is given the same task as SafeMAP in a single request: the project's C sources and headers, the planned public signatures (C names kept), and the Rust struct definitions those signatures use, with private helpers allowed. It receives none of SafeMAP's analysis facts, plans, rules, or repair. Its output is scored with the same per-unit harness after making top-level functions public. Giving the interfaces isolates what SafeMAP adds beyond interface planning. The *plain LLM* baseline (amendment 2) measures the planning itself: the same request without the planned signatures and struct definitions, so the model sees only the C sources and the safety rules. Its output is scored the same way; a function with any other signature cannot be called by the harness and is recorded as `interface_mismatch`.

4 Evaluation methodology

4.1 Earlier evaluation

A first frozen evaluation (protocol v1) ran SafeMAP on inih, cJSON, and libcsv and accepted 0 of 142 units. Every one of six defects independently forced that result: the signature planner mapped C types incorrectly; rules fired on partial matches and emitted template code; struct definitions in headers were not collected; acceptance was all-or-nothing per project; the C2Rust setup shadowed libc headers; and the harness covered only scalar shapes. Because the v1 corpus was inspected while diagnosing these defects, it is used here only as development data, and the corrected system was evaluated on newly selected libraries under a new protocol.

4.2 Protocol

The v2 protocol fixed the selection rule, configurations, model, acceptance policy, and reporting before any SafeMAP execution on the held-out corpus. A library qualifies if it is a public Git repository with an OSI-approved license, its production code is one `.c` file plus headers that compiles against libc alone, the file has 150–3,500 lines, and it was not used during development. All functions in the file are retained. The protocol and corpus manifest were sealed and the implementation frozen at commit `62cf430` before the corpus was materialized; a materialization script refuses to run on any other commit. The full test suite passed at the freeze.

Four libraries were selected first: parson, sds, tinyexpr, and microtar. After their results were seen, we judged four libraries too few and added six more by the same rule

Table 2 Held-out corpus: ten libraries from two selection rounds.

Library	Round	C LOC	Functions	Units	Ptr. params	Complexity
microtar	Original	376	25	25	39	68
parson	Original	2494	144	138	206	633
sds	Original	1328	45	45	27	292
tinyexpr	Original	1025	48	37	28	306
argparse	Extension	395	13	13	27	101
json-parser	Extension	1186	10	10	15	227
printf	Extension	914	34	20	27	296
semver	Extension	638	33	33	29	127
tiny-AES-c	Extension	572	20	20	33	47
tiny-regex-c	Extension	528	18	15	23	126
Total		9456	390	356	454	2223

(amendment 1): json-parser, tiny-regex-c, tiny-AES-c, semver.c, argparse, and printf. Nine candidates qualified in a fresh screen; six were chosen for domain diversity before any SafeMAP execution on them. The implementation, configuration, and modes were unchanged. We report the two rounds separately and pooled. After all results were seen we added one baseline (amendment 2), the plain LLM described above, run three times with the frozen implementation by a script that refuses to run if the implementation or configuration differs from the freeze. Table 2 characterizes the corpus.

4.3 Configuration

All LLM configurations, SafeMAP and baseline alike, use Google's `gemini-3.5-flash-lite` at temperature zero through its OpenAI-compatible API, with at most three repairs per unit, 1,000 generated inputs per unit, and seed 20260728. The evaluation used the provider's free tier. A request still refused after the client's retries stops the run rather than being scored as a failure; the run resumes after the quota resets, reusing every finished (project, configuration) result. Because LLM output can vary even at temperature zero (Ouyang et al. 2025), the LLM configurations were run three times on every held-out library with identical inputs; the first run is the primary result. Tool versions: rustc and Clippy 1.99.1, Clang 23.1.3; C2Rust 0.22.1 with LLVM 14 in a pinned container.

4.4 Development data

Development data comprises 40 authored microbenchmarks, six authored modules, ten LLVM test-suite programs pinned at commit `6cdc54e` (LLVM Project 2026), and the three v1 libraries. Their results informed the corrections, so they are not evidence of generalization. The component ablations are run on the v1 libraries with the frozen implementation and are reported as development data, which keeps the held-out LLM run within the provider's quota.

Table 3 SafeMAP (guided) outcome cascade per library (first run). *Crate*: the unit is in a crate that passed `cargo check`, tests, and Clippy.

Library	Units	Candidate	Generated	Crate	Policy	Accepted
microtar	25	3	2	2	2	2
parson	138	72	4	4	4	3
sds	45	30	4	4	4	2
tinyexpr	37	32	24	0	0	0
argparse	13	0	0	0	0	0
json-parser	10	3	2	2	2	1
printf	20	18	0	0	0	0
semver	33	17	4	4	4	3
tiny-AES-c	20	1	1	1	1	1
tiny-regex-c	15	7	6	6	6	4
Total	356	183	47	23	23	16

4.5 Mutation analysis

To test whether the per-unit check rejects wrong code, we used mutation analysis (Jia and Harman 2011). We applied single-operator mutations (relational, arithmetic, logical, boolean, and integer-literal changes) to every accepted function of SafeMAP's held-out runs and of its v1-library run, kept the mutants that compile, and re-ran each unit's differential test with 300 inputs from a different seed. We classified every surviving mutant by hand as equivalent (no input can distinguish it), an input gap (a distinguishing input exists but was not generated), or an observation gap (the harness does not compare the differing output).

5 Results

5.1 RQ1: Accepted units

Table 3 follows each held-out unit through the pipeline. In the first run, of 356 units, the screen marked 183 as candidates, SafeMAP generated Rust for 47, and accepted 16 (4.5% of all units; 8.7% of candidates). The original four libraries contributed 7 accepted units and the extension 9. Accepted units are small leaf functions: a string hash, hexadecimal and UTF-8 helpers, rounding, error strings, version-string checks, regular-expression character classes, and the AES `xtime` step.

Of the 31 generated units not accepted, 24 are from tinyexpr. They compiled, but the crate failed Clippy: two translated constants, `pi` and `e`, keep C's literal digits, which Clippy's default-deny `approx_constant` lint rejects. Because Clippy runs once per crate, this rejected every unit in the crate. We report the outcome as frozen; it is a limitation of crate-level checks within a per-unit policy (Section 7).

Post-hoc sensitivity. To measure that limitation, we re-scored the frozen tinyexpr crate with only `approx_constant` allowed, using the unchanged crate checks and differential harness (`scripts/posthoc_lint_sensitivity.py`). The crate then passed compilation and Clippy, and 23 of the 24 units matched the C functions on

Table 4 Recorded reasons for functions not accepted by SafeMAP (guided), all ten held-out libraries (first run).

Reason	Functions
Manual refactor (global state, I/O, allocation protocols)	159
No safe, testable interface for a C type	116
Requires unsafe (raw memory, casts)	36
Crate-level Clippy check failed	24
Unsupported construct (union, function pointer, asm)	12
C undefined on most test inputs	4
Differential test failed	3
Callee not translated	3
No synthesis plan	3

every defined input; the remaining unit, `negate_logical_not`, disagreed with C. Per-unit lint checking would therefore have accepted 39 units in the first run instead of 16. This analysis was not part of the sealed protocol and the primary result remains 16.

5.2 RQ2: Why units are not accepted

Table 4 lists the recorded reasons. Most units never reach translation. 159 functions require manual refactoring, mainly because they read or write mutable global state, perform I/O, or follow allocation protocols that cross function boundaries. 116 use a C type for which the planner finds no safe, testable interface: in parson and sds the central types (`JSON_Value *`, `sds`) are opaque handles or point into the middle of an allocation, and `void *` and `char **` parameters are common. 36 require unsafe operations such as pointer arithmetic beyond a known length or casts between pointer types. Only a small fraction of losses occur after translation: 3 differential-test failures, and 4 units whose C oracle had undefined behavior on most generated inputs. Two of these show a limit of signature-level planning: sds's `sdsll2str` writes up to 21 bytes into a caller buffer, a requirement its signature does not express, so every generated buffer overflowed. The other two overflow signed arithmetic on most inputs.

5.3 RQ3: Baselines

Table 5 compares the configurations. SafeMAP accepted 16 units in the first run, the one-shot LLM given the planned interfaces 9, the plain LLM 1, rules alone 1, and C2Rust none. Over three runs SafeMAP and the interface-given baseline overlap (Table 8), so we do not claim that SafeMAP's loop outperforms it. The plain LLM shows what the planned interfaces contribute: for four libraries its output used `unsafe` code or raw pointers and was rejected outright, and in the other six, of 43 planned functions, 24 were absent from the output and 12 had other signatures, so they could not be checked against the C function; 6 were in crates that failed compilation or Clippy, and 1, in semver, was accepted. SafeMAP used 40 LLM requests and the one-shot baseline 8 (one per library with planned interfaces). The one-shot baseline benefits from SafeMAP's interface planning: without the planned signatures, its output could not be tested function by function. Raw C2Rust output compiled for seven

Table 5 Accepted units per configuration on the held-out libraries (first run). One-shot: given the planned signatures; plain: C sources only (amendment 2). LLM calls: guided / one-shot. †C2Rust output did not compile.

Library	Units	Guided	One-shot	Plain	Rules	C2Rust	LLM calls
microtar	25	2	1	0	1	0	1 / 1
parson	138	3	4	0	0	0†	4 / 1
sds	45	2	0	0	0	0†	4 / 1
tinyexpr	37	0	0	0	0	0	17 / 1
argparse	13	0	0	0	0	0	0 / 0
json-parser	10	1	1	0	0	0	2 / 1
printf	20	0	0	0	0	0†	0 / 0
semver	33	3	0	1	0	0	5 / 1
tiny-AES-c	20	1	1	0	0	0	1 / 1
tiny-regex-c	15	4	2	0	0	0	6 / 1
Total	356	16	9	1	1	0	40 / 8

Table 6 Development data: accepted units per configuration.

Corpus	Units	Guided	One-shot LLM	Rules only	C2Rust
Microbenchmarks	81	72	35	35	0
Authored modules	25	24	18	18	0
LLVM programs	32	10	8	6	0
v1 libraries	142	22	21	0	0

of ten libraries; for parson, sds, and printf it failed on the pinned toolchain because its `va_list` handling and packed-struct field references are rejected by current rustc. Where it compiled, it contained unsafe code and raw-pointer interfaces throughout, so no unit met the policy.

Development data show the same pattern (Table 6): on the v1 libraries the two accepted 22 and 21 units. Only on the authored microbenchmarks, which the rules were written for, is SafeMAP clearly ahead (72 of 81 units against 35).

5.4 RQ4: Component ablations

Table 7 reports the ablations on the v1 libraries. The full system accepted 22 of 142 units. Removing safe-signature planning reduced this to 2, while raising LLM requests from 23 to 126, because the model must then invent interfaces that rarely compile with their callers. Removing idiom plans reduced acceptance to 3. Withholding analysis facts and plans from the prompt or removing compiler-feedback repair changed nothing (22), removing dependency grouping lost one unit (21), and removing pointer roles lost one (21). The planning that determines *what* interface each function gets, rather than the guidance given to the LLM, accounts for the accepted translations.

Table 7 Component ablations on the v1 libraries (development data): accepted / all units.

Configuration	cJSON	inih	libcsv	LLM calls
Full system	15/109	1/10	6/23	23
Facts and plan withheld from prompt	15/109	1/10	6/23	22
Without pointer roles	15/109	0/10	6/23	23
Without safe signatures	2/109	0/10	0/23	126
Without dependency grouping	14/113	1/10	6/23	21
Without idiom plans	3/109	0/10	0/23	3
Without compiler-feedback repair	15/109	1/10	6/23	22

Table 8 Accepted units across repeated LLM runs (identical inputs and configuration). One-shot: given the planned signatures; plain: C sources only.

Run	Libraries	Guided	One-shot	Plain
Run 1	10	16	9	1
Run 2	10	15	15	1
Run 3	10	12	14	2

5.5 RQ5: Reliability of acceptance

Mutation analysis. Of 152 compiling mutants of accepted functions (all ten held-out libraries and the v1 libraries), the per-unit check rejected 135 (88.8%). Of the 17 survivors, 4 are equivalent: `>` replaced by `>=` directly after an equality test, the same change where both branches yield the same value on ties, a loop bound that a terminator written just before already guarantees, and an unreachable `else`. Excluding them, the check rejected 135 of 148 non-equivalent mutants (91.2%). Twelve survivors are input gaps: they change behavior only on a few specific values, such as one or two digits in sds's hexadecimal converter (six mutants), a string of exactly 255 or 256 bytes, or one error code, which uniformly generated inputs rarely hit. One is an observation gap: a copy routine's output differs only after the string terminator, which the harness, like a C caller, does not read. Targeted inputs from the C function's literals and comparison bounds would close most input gaps.

Repeated runs. Table 8 reports accepted units in three runs of the LLM configurations with identical inputs. SafeMAP accepted 16, 15, and 12 units, the one-shot baseline 9, 15, and 14, and the plain LLM 1, 1, and 2. Eight of the ten libraries gave the same SafeMAP result in every run. Most of the variation comes from semver: in one run of each configuration, one generated function failed a crate-level check and so rejected all four planned units of that crate, the same mechanism as tinyexpr. The accepted sets are therefore stable at the function level but sensitive to crate-level failures.

6 Discussion

Coverage is the limiting factor. The per-function acceptance check is feasible on production C and rejects most wrong code, but few functions reach it. Table 4 is the main finding: the obstacles are global state, opaque and aliasing pointer types, and raw memory manipulation. These are API design questions that a function-local translator cannot answer alone; a migration would need to change the library's interface, for example by turning global state into an explicit context or an `sds` string into an owned type, before SafeMAP's checks apply.

Interfaces matter more than prompts. On the held-out libraries the same model accepted 1–2 units from the C code alone and 9–15 when told the planned interfaces, within the run-to-run range of SafeMAP's guided loop. On the development libraries, removing signature planning from the guided loop reduces it from 22 to 2 units. Withholding analysis facts from the prompt changed nothing on the v1 libraries. This suggests that LLM translators benefit most from being told the target interface, which is consistent with CRUST-Bench's use of hand-written interfaces (Khatry et al. 2025).

Refusal as output. Every refused unit carries a reason and the plan that produced it. For a migration team this identifies which functions can be migrated mechanically and which need design work.

7 Threats to validity

Sample. Ten small libraries with one production file each do not represent large applications or multi-file projects. The extension was decided after the first four results were seen; it used the same rule and frozen code, and we report both rounds separately.

Model and variance. One low-cost model was used, and three runs are too few to estimate run-to-run variance precisely; the guided and one-shot ranges overlap, so their difference is not established. The plain LLM baseline was added after the other results were seen. A stronger model may translate more of the planned units, although it cannot change the units refused before translation, which are most of the losses.

Oracle. Differential testing on generated inputs is evidence, not proof. Mutation analysis shows that 8.8% of non-equivalent compiling mutants survive, mainly where only a few specific inputs reach a branch. Excluding inputs with undefined C behavior removes inputs the sanitizers detect, not all undefined behavior.

Crate-level checks. Clippy runs on the whole crate, so one lint error rejects every unit in it (tinyexpr, and semver in one repeated run). This makes the reported numbers conservative; the post-hoc analysis in Section 5 puts the cost at 23 units in the first run.

No eligibility labels. We did not collect independent labels of which functions could be migrated safely, so we make no claim about the precision or recall of the screen; refusals are reported as recorded.

Toolchain. C2Rust 0.22.1's output does not compile with the pinned nightly on three libraries; another toolchain might compile it, but it would still use unsafe code and fail the policy.

Development data. The development corpora and the v1 libraries shaped the implementation; results on them are not evidence of generalization, and the ablations are run on them.

8 Conclusion

SafeMAP translates C functions to safe Rust only when it finds a safe interface and accepts them only when they match the original C under sanitizers, so every accepted function carries its own evidence of equivalence. On ten held-out libraries it accepted 12–16 of 356 units across three runs. The same LLM accepted 1–2 units from the C code alone and 9–15 when given SafeMAP's planned interfaces; rules alone accepted 1 and raw C2Rust 0. The interface planning, not the LLM guidance, is what makes these translations possible to produce and check. Most units are refused before translation because of global state, opaque or aliasing pointer types, or unsafe memory operations. Per-function, oracle-checked migration is therefore feasible but narrow: its reach is set by how much of a C library's interface already has a safe Rust counterpart. Two limits of the acceptance check itself are measurable and fixable: crate-level lint checks cost 23 units in the first run, and most surviving mutants need inputs that target specific values. Future work should run lint checks per unit, generate inputs from the C function's literals and bounds, and plan interface changes such as explicit contexts for global state.

Statements and Declarations

Funding

No funding was received for conducting this study or preparing this manuscript.

Competing interests

The authors declare no competing financial or non-financial interests relevant to the content of this article.

Data availability

The source code, evaluation protocol, corpus manifests, configuration, frozen implementation record, raw results, and scripts that generate every table are archived on Zenodo (https://doi.org/10.5281/zenodo.23494717) and available in the SafeMAP repository (https://github.com/jscyril/SafeMAP); the implementation evaluated here is commit `62cf430`. The result rows of every run are in `reports/publication/v2`, and `scripts/generate_v2_manuscript_tables.py` regenerates every table from them. The held-out libraries are public repositories pinned to the revisions recorded in the manifests.

Ethics approval

Not applicable. The study involved no human participants.

Use of generative AI

Generative AI tools were used in this work. ChatGPT assisted with generating the synthetic development benchmarks. ChatGPT and Claude assisted with implementing and testing the software, writing the evaluation scripts, drafting and revising the manuscript, and refining its language. The evaluated pipeline itself calls Gemini as described in the paper. The authors reviewed and validated all AI-assisted code, data, and text, and remain responsible for the accuracy, originality, and integrity of the manuscript.

Author contributions

Jacob Sebastian Cyril: conceptualization, methodology, software, investigation, formal analysis, visualization, and writing—original draft. Benita Jose Chalissery: research guidance and manuscript review. NewBegin M: supervision and manuscript review. All authors approved the final manuscript.